

SEONGLAE CHO

London, UK | sungle3737@gmail.com | +44 7356-161248 | [LinkedIn](#) | [GitHub](#) | [Website](#) | [Scholar](#)

Senior AI Engineer with 4 years of full-stack software engineering and a first-author ICML 2026 main publication. Builds ML systems, agent infrastructure, and developer tooling, from research to production.

EXPERIENCE

- HOLISTIC AI**, London, UK · *Senior AI Engineer* May 2025 – Present
- Built a trace-based AI agent evaluation framework with structured traces, quality criteria, and graph-based monitoring.
 - Built an FSM inference system reconstructing agent finite-state machines from execution traces for 95%+ behavior-coverage analysis.
 - Integrated ReBAC access control into the AI Governance platform with database-level policy enforcement.
 - Integrated agent features into an event-driven architecture (EDA) via an AWS Lambda-triggered Event Bus.
 - Promoted to Senior within one year of joining.
- KAKAO MOBILITY**, Seoul, South Korea · *Software Engineer, Digital Twin Team* Dec 2021 – Sep 2022
- Led a 3-engineer team shipping a vector-annotation platform processing 10K+ daily annotations for the autonomous-driving pipeline.
 - Built Kubernetes infrastructure with Grafana alerting and isolated namespaces, cutting incident response time 60%.
 - Rewrote a C++ 3D-projection algorithm in Rust with Node.js bindings for a 40% performance gain.
- STRYX**, Seoul, South Korea · *Software Engineer, Geospatial AI Startup* Nov 2019 – Dec 2021
- Cut build time 70% by consolidating 5 repositories into a unified monorepo.
 - Reduced the Docker image 90% (2GB to 180MB) via multi-stage CI builds, accelerating deployment 3x.
- YONSEI UNIVERSITY, Data & Language Intelligence Lab**, Seoul, South Korea · *Research Intern, Prof. Dongha Lee* Mar 2023 – Jan 2024
- Interpretable-summarization research on LLMs; co-authored first publication (RTSUM, NAACL 2024 Demo).

PROJECTS

- Vect** · local-first multi-agent coding IDE, ~80k LOC solo; orchestrates Claude Code and 4 other coding-CLI agents on a CRDT-merged workspace via MCP. Mar 2026 – Present
- MCP-Notion** · Model Context Protocol SSE server searching Notion pages and converting them to Markdown for agent interoperability. Jan 2025 – Feb 2025
- Conference Arena** (confarena.com) · paper-rating leaderboard web app with Elo-style ranking. Jan 2026 – Feb 2026
- TPCxAI** · scalable big-data pipelines on a Hadoop + Spark cluster in GCP for end-to-end AI workloads. Mar 2024 – May 2024
- ReSRer** · pipeline indexing 21M documents into Milvus, accelerating LLM training across a 4xRTX 3090 cluster. Sep 2023 – Jan 2024
- LLaMa2GPTQ** · local LLM with 4-bit GPTQ quantization (~75% memory) and RAG vector search. Jun 2023 – Jul 2023
- Texonom** · Zettelkasten knowledge system; ANN retrieval over 30,000 pages in Postgres pgvector. Nov 2021 – Jun 2023

PUBLICATIONS

- Cho, S., Wu, Z., & Koshiyama, A. (2025). [CorrSteer](#): Generation-Time LLM Steering via Correlated Sparse Autoencoder Features. **ICML 2026 main**.
- Cho, S., Wu, Z., & Koshiyama, A. (2026). [Control Reinforcement Learning](#): Interpretable Token-Level Steering of LLMs via Sparse Autoencoder Features. ICLR 2026 Workshop on Principled Design for Trustworthy AI.
- Cho, S., Wu, Z. et al. (2026). [The Confidence Manifold](#): Geometric Structure of Correctness Representations in Language Models. ICLR 2026 Workshop on Agentic AI in the Wild.
- Cho, S. et al. (2026). PaaT (Probe as a Tool): self-regulating agents via a callable activation probe. ICML 2026 FAGen Workshop.
- Cho, S. et al. (2026). AgentRoom: Concurrent Multi-Agent Coding in a CRDT-Backed Shared Workspace. ICML 2026 FAGen Workshop.
- Cho, S., Wu, Z. et al. (2026). Automata from Agent Traces: Failure and Next-Step Prediction. ICML 2026 AIWILD Workshop.
- Cho, S., Wu, Z. et al. (2026). Are Single-Token Sparse Autoencoder Features Irreplaceable? The Impact of Training Methodology. EMNLP 2026 under review.
- Cho, S. et al. (2025). [FaithfulSAE](#): Towards Capturing Faithful Features with Sparse Autoencoders without External Datasets Dependency. ACL 2025 SRW.
- Wu, Z., Cho, S. et al. (2025). AgentGraph: Trace-to-Graph Platform for Interactive Analysis and Robustness Testing in Agentic AI Systems. AAAI 2026 Demo.
- Wu, Z., Cho, S.* et al. (2025). [LibVulnWatch](#): A Deep Assessment Agent System and Leaderboard for Uncovering Hidden Vulnerabilities in Open-Source AI Libraries. ICML 2025 TAIG Workshop.
- Cho, S. et al. (2024). [RTSUM](#): Relation Triple-based Interpretable Summarization with Multi-level Saliency Visualization. NAACL 2024 System Demonstrations.

EDUCATION

- University College London** · *MSc in AI for Sustainable Development, Distinction* Sep 2024 – Sep 2025
- Dissertation: Representation-Level Steering Improves Task Performance and Safety in LLMs. Supervisor: Prof. Philip Treleaven; Mentor: Zekun Wu.
- Yonsei University** · *BE in Computer Science* Mar 2017 – Aug 2024

SKILLS

- AI**: Python, PyTorch, AI agents, Langfuse, LangSmith, DDP/FSDP, ONNX, PydanticAI, RAG, LangGraph, Agent SDK.
- Service**: Full-stack, TypeScript, Node.js, PostgreSQL, Redis, Rust, Faiss, Milvus, Vite, React, Next.js.
- Infra**: Kubernetes, Linux, Git, CI/CD, Docker, EDA, Apache Spark, Ansible, ETL, AWS, GCP, Hadoop, Terraform, SST.

AWARDS

- HERMES, 1st Place (£3,000), Holistic AI Hackathon** (Nov 2024) · cut stereotypical generation from 90% to 70% with an SAE-derived steering vector.

- **MBTIGPT, 1st Place (₩3,000,000), Yonsei GenAI Competition** (2023) · AI service with 1,000+ users, RAG over chat history with Redis and Faiss.

SERVICE

- Reviewer, ICML 2026 (Silver Reviewer Award) · Reviewer, ACL Rolling Review (March 2026).
- Mandatory military service completed as an Industrial Technical Personnel (alternative service).